

# Module 2

## Efficient Fine-Tuning

Doing more with less



# Learning Objectives

## By the end of this module you will:


- Understand what fine-tuning is and why we do it
- Learn what parameter-efficient fine-tuning is and what the popular strategies are
- Understand the limitations of parameter-efficient fine-tuning
- Gain knowledge about data preparation best practices



# Fine tuning vs. transfer learning

They are often referenced interchangeably

- Transfer learning
  - Apply a general pre-trained model to a new, but related task
- Fine tuning
  - Use a general pre-trained model and then train that model further
- Transfer learning  $\sim$  fine tuning
  - Train it more
  - Train on different data



# How to leverage a pre-trained foundation model?




Examples:

- T5
- BloombergGPT
- GPT-4



# How to leverage a pre-trained foundation model?


Pre-trained



Examples:

- T5
- BloombergGPT
- GPT-4

Feature  
Extraction




Example:

Use BERT embeddings as inputs  
to a random forest classifier  
→ classify movie reviews

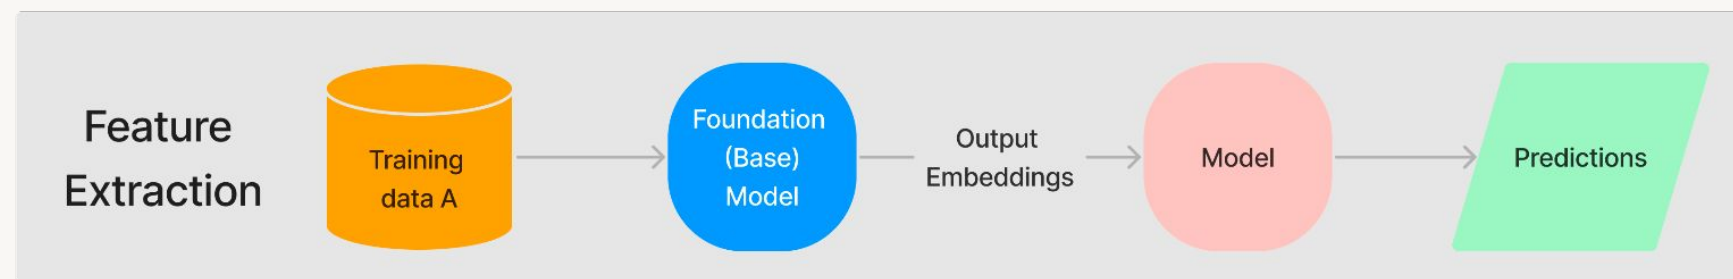


# How to leverage a pre-trained foundation model?



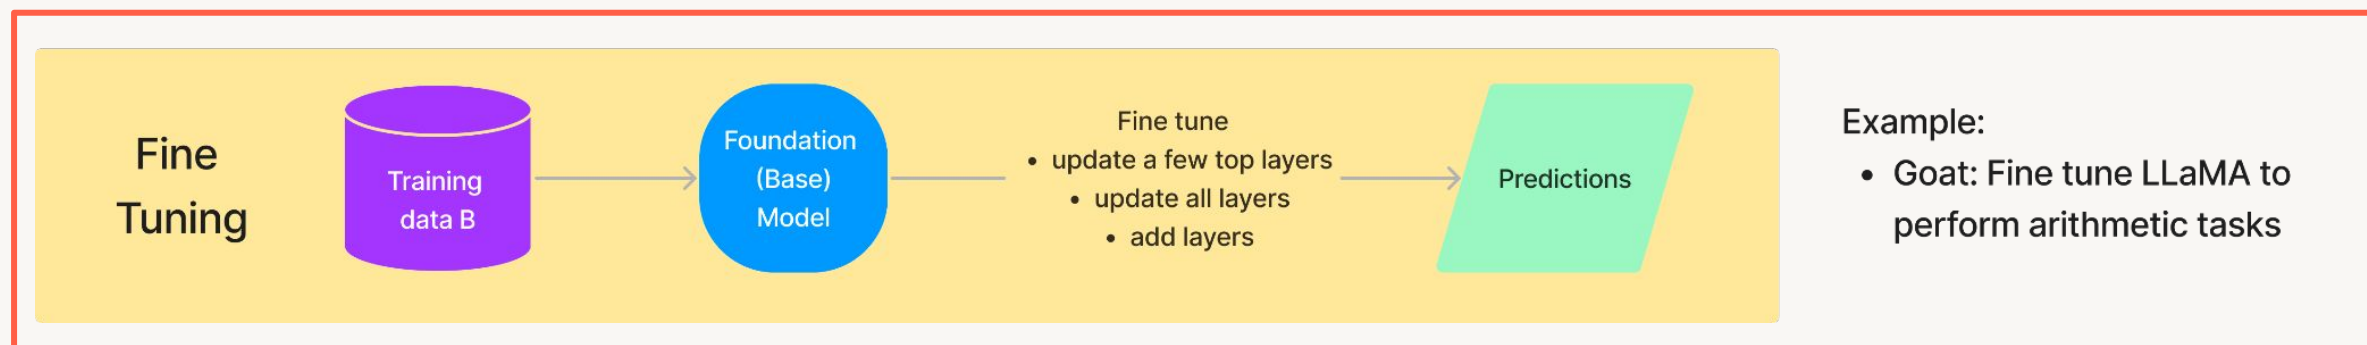
Examples:

- T5
- BloombergGPT
- GPT-4



Example:

Use BERT embeddings as inputs to a random forest classifier  
→ classify movie reviews



Example:


- Goat: Fine tune LLaMA to perform arithmetic tasks



# Why fine tuning?

Leverage an effective pre-trained model on our own data – it's *not* new

- Improve performance downstream
  - Different pre-trained vs fine-tuned tasks
  - Different domains
- Ensure regulatory compliance
- Not new:
  - [ULMfit paper](#) in 2018

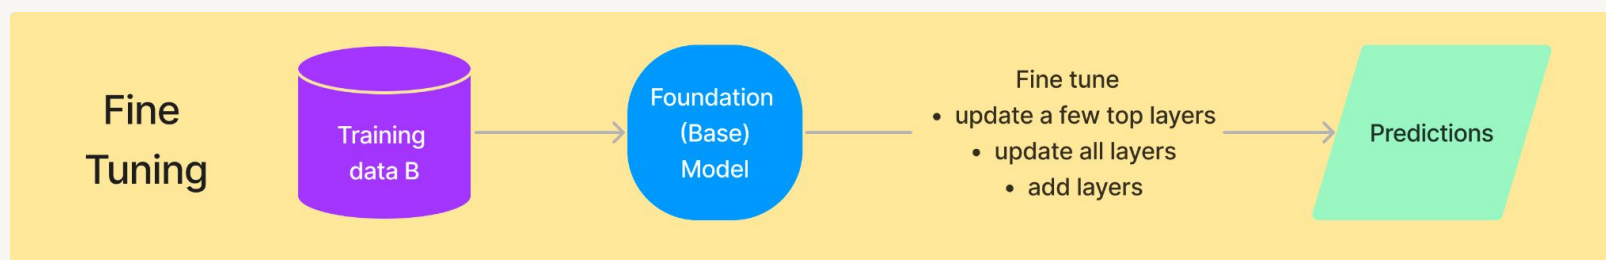


Source: [Howard and Ruder 2018](#)

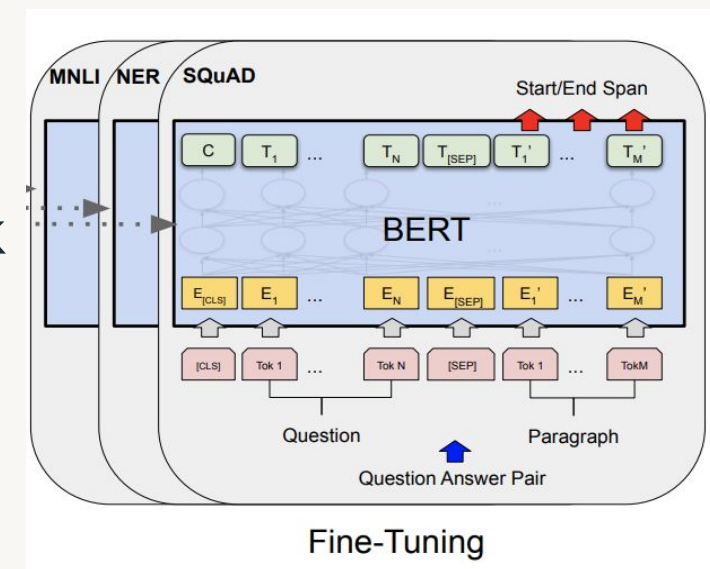


# Fine tune = update foundation model weights

AKA parameter fine tuning



- Update more layers = better model performance
- Full fine-tuning typically produces one model per task
  - Serve one model per task
  - May forget other pre-trained tasks: catastrophic forgetting
- Full fine-tuning LLMs is expensive. How to avoid it?
  - X-shot learning
  - Parameter-efficient fine tuning



Source: [Devlin et al 2019](#)





# X-shot learning

Provide several examples of new tasks

Prompt engineering

= developing prompts

= prompt design

= *hard/discrete prompt tuning*

**Not updating model weights**

```
pipeline(  
    """For each tweet, describe its sentiment:  
  
    [Tweet]: "I hate it when my phone battery dies."  
    [Sentiment]: Negative  
    ###  
    [Tweet]: "My day has been 👍"  
    [Sentiment]: Positive  
    ###  
    [Tweet]: "This is the link to the article"  
    [Sentiment]: Neutral  
    ###  
    [Tweet]: "This new music video was incredible"  
    [Sentiment]: """)
```

Instruction

Few-shot  
examples

Prompt



# Pros and cons of X-shot learning

Also known as in-context learning

## Pros

- No need for huge labeled training data
- No need to create a copy of model for each task
  - Simplify model serving
- Text prompts feel interpretable

## Cons

- Manual prompt engineering
- Prompts are specific to models
- Context length limitation
  - Add more examples? Less space for instructions
  - Longer context = higher latency
  - LLMs forget middle portion
    - [Liu et al 2023](#) (released in July)
    - Longer context window is not the solution!
- **Performance might still be lackluster**



# Fine-tuning outperforms X-shot learning

## Example: G0od at Arithmetic Tasks (Goat-7B)

- Foundation model = Llama
- Trained on 1M synthetic data samples
- Accuracy outperforms
  - Few-shot PaLM-**540B** (*much bigger model !!* )
  - GPT-4
    - Typically doesn't do well in arithmetics; accuracy ~0
- SOTA on arithmetic benchmark (BIG-bench)
- Supervised **instruction fine-tuning**
  - Trained using LoRA on a 24GB VRAM GPU
    - LoRA to be covered soon!



[Image Source](#)

| Goat-7B   |
|---|
| <b>3978640188 + 42886272 =</b><br>3978640188 + 42886272 = 4021526460  |
| <b>4523646 minus 67453156</b><br>4523646 - 67453156 = -62929510   |
| <b>Calculate 397 x 4429</b><br>397 × 4429<br>= 4429 × (300 + 90 + 7)<br>= 4429 × 300 + 4429 × 90 + 4429 × 7<br>= 1328700 + 398610 + 31003<br>= 1727310 + 31003<br>= 1758313 |
| <b>What is 8914/64?</b><br>8914 - 64 × 100 = 8914 - 6400 = 2514<br>2514 - 64 × 30 = 2514 - 1920 = 594<br>594 - 64 × 9 = 594 - 576 = 18<br>Therefore, 8914 ÷ 64 = 139 R 18   |

Source: [Liu and Low 2023](#)  
(released in May)



# Important observations about Goat

- Instruction fine-tuned
- Multi-task serving

## Task 1 Addition

### Goat-7B

**3978640188 + 42886272 =**  
 $3978640188 + 42886272 = 4021526460$

**4523646 minus 67453156**  
 $4523646 - 67453156 = -62929510$

## Task 3 Multiplication, mix of NL + mathematical symbols

**Calculate 397 x 4429**  
 $397 \times 4429$   
 $= 4429 \times (300 + 90 + 7)$   
 $= 4429 \times 300 + 4429 \times 90 + 4429 \times 7$   
 $= 1328700 + 398610 + 31003$   
 $= 1727310 + 31003$   
 $= 1758313$

## Task 2 Subtraction, using natural language (NL)

**What is 8914/64?**  
 $8914 - 64 \times 100 = 8914 - 6400 = 2514$   
 $2514 - 64 \times 30 = 2514 - 1920 = 594$   
 $594 - 64 \times 9 = 594 - 576 = 18$   
Therefore,  $8914 \div 64 = 139 \text{ R } 18$

## Task 4 Division, mix of NL + mathematical symbols



# Instruction-tuned, multi-task LLM

Instruction-tuned = tune general purpose LLMs to follow instructions

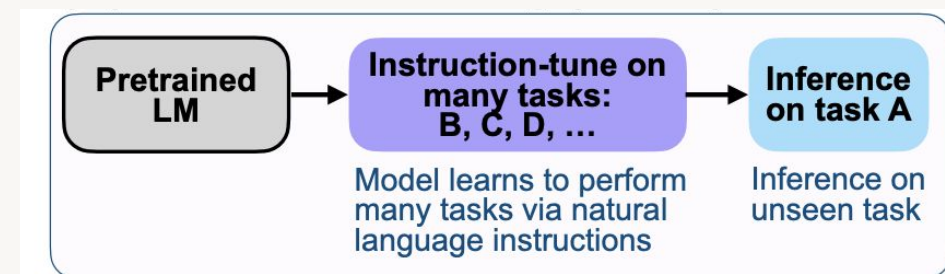
## FLAN (Fine-tuned LAnguage Net)

- Foundation model = 137B model
- Instruction-tuned on over 60 NLP datasets with different task types
  - Task types: Q/A, translation, reasoning, comprehension, etc.
- Examples
  - T5 -> FLAN-T5
  - PaLM -> FLAN-PaLM

## Dolly



- Foundation model = Pythia-12B
- Instruction-tuned on 15k prompt/response pairs
  - Task types: Q/A, classification, information extraction, etc.



Source: [Wei et al 2022](#)



# Quick recap

We want efficient training, serving, and storage

- Full fine-tuning can be computationally prohibitive
  - Memory usage: activation, optimizer states, gradients, parameters
  - This gives the best performance
- Compromise: Do *some*, but not full, fine-tuning
  - Saves cost to use low-memory GPUs
- We want multi-task serving, rather than one model per task
  - E.g. one model for Q/A, summarization, classification

•  
*Enter **parameter-efficient** fine-tuning*



# Parameter-efficient fine-tuning (PEFT)



# 3 categories of PEFT methods

## Additive

- Soft prompt
  - Prompt tuning
  - Prefix tuning

## Selective

- Akin to updating a few foundation model layers
  - [BitFit](#)
    - Only updates bias parameters
  - [Diff Pruning](#)
    - Creates task-specific “diff” vectors and only updates them

## Re-parameterization

- Decompose weight matrix updates into smaller-rank matrices
  - LoRA





# We will cover additive and reparameterization

## Additive

- Soft prompt
  - Prompt tuning
  - Prefix tuning

## Selective

- **Model quality performance is not as good**
- Akin to updating a few foundation model layers
  - BitFit
    - Only updates bias parameters
  - DiffPruning
    - Creates task-specific “diff” vectors and only update them

## Re-parameterization

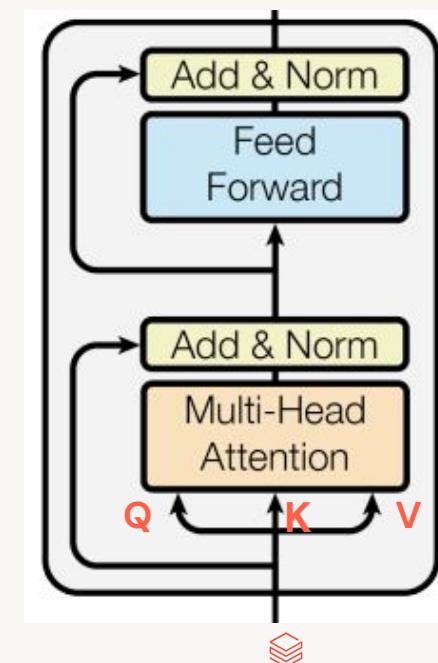
- Decompose weight matrix into smaller-rank matrices
  - LoRA



# High-level overview of PEFT

Active research area: >100 papers in last few years!

- Additive: Add new tunable layers to model
  - Keep the foundation model weights frozen and update only the new layer weights
- Reparameterization: Decompose a weight matrix into lower-rank matrices
- Implementation:
  - Acts on the core Transformer block
    - Basic multi-head attention and/or feed forward network
  - Some act specifically on the weight matrices: **Query, Key, Value**
    - These matrices pass information from one token to another



Source: [Vasmani et al 2021](#)

# Additive: Prompt Tuning (and prefix tuning)

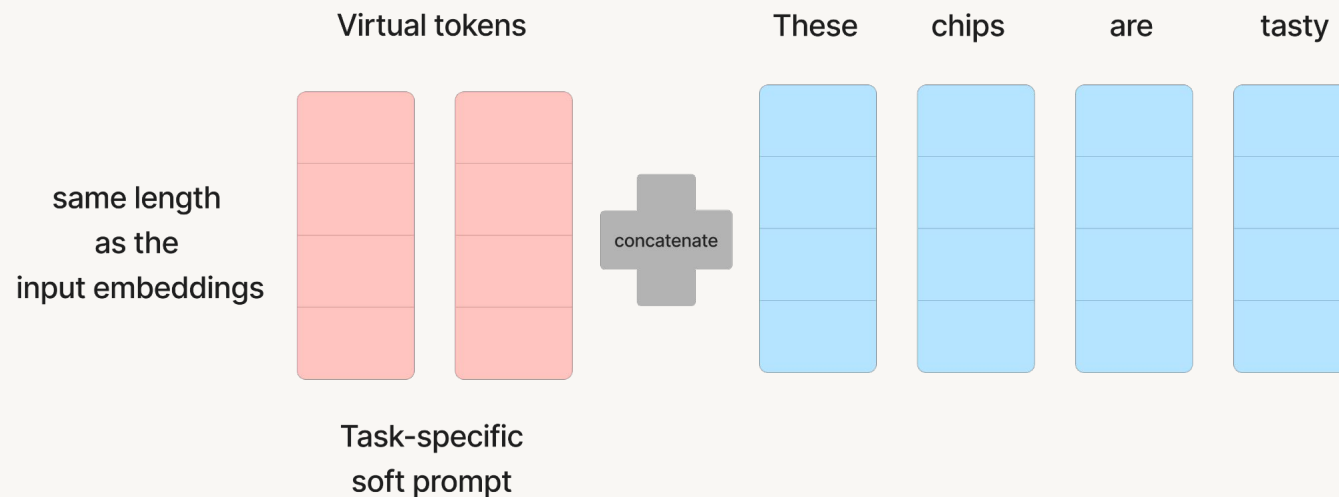


# Soft prompt tuning

Concatenates trainable parameters with the input embeddings

- Learn a new sequence of task-specific embeddings
- We call this prompt tuning, not model tuning, because we only update prompt weights

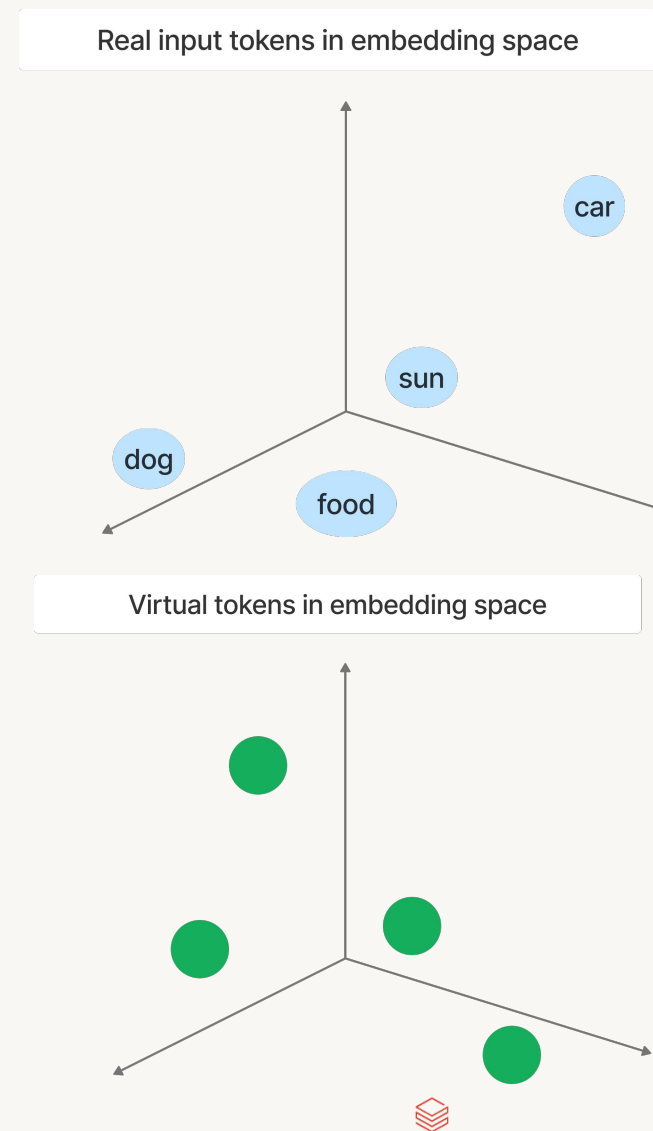
Task: Classify product reviews; task batch = 1



# What are these *virtual* tokens?

Goal: remove manual element of engineering prompts!

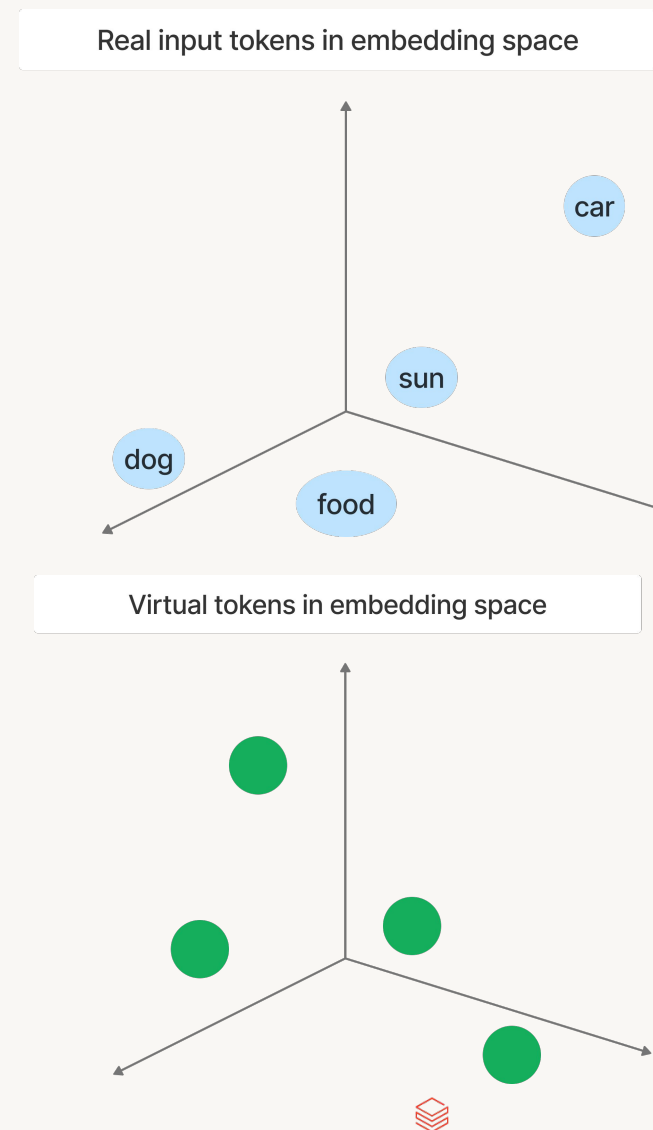
- Randomly initialized embedding vectors
- Not part of vocabulary
- Analogy:
  - Bitcoin: We can't touch it like cash. We don't know how it "looks", but it exists and works.



# What are these *virtual* tokens?

Goal: remove manual element of engineering prompts!

- Randomly initialized embedding vectors
  - We can also initialize to discrete prompts
  - But random initialization is nearly as good as informed initialization ([Qin and Eisner 2021](#))
- Not part of vocabulary
- Analogy:
  - Bitcoin: We can't touch it like cash. We don't know how it "looks", but it exists and works.

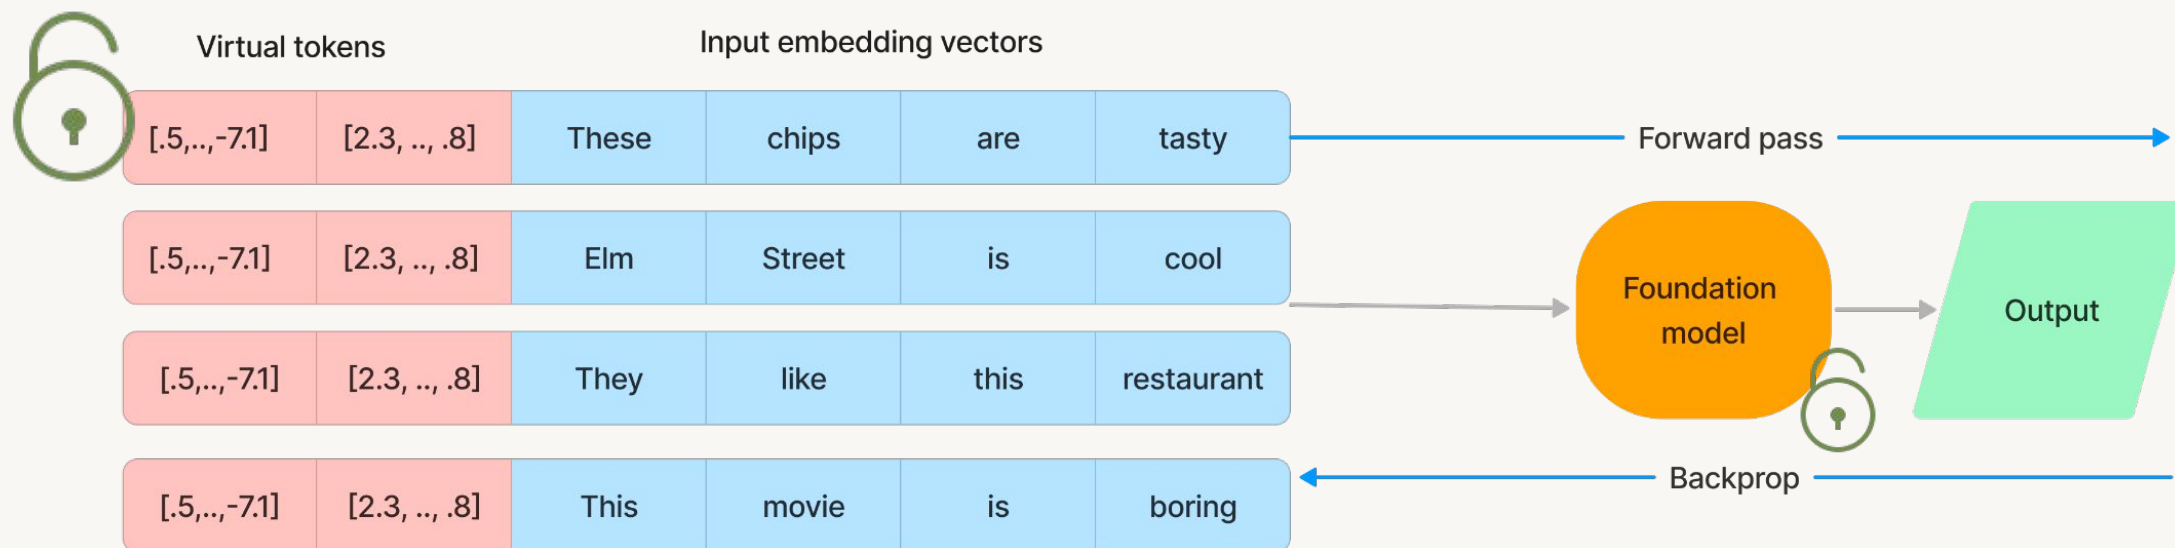


# Compare full fine-tuning vs prompt tuning

Scenario: full fine-tuning

Backprop: update **all** weights based on loss

Task: Classify sentiment; task batch = 4



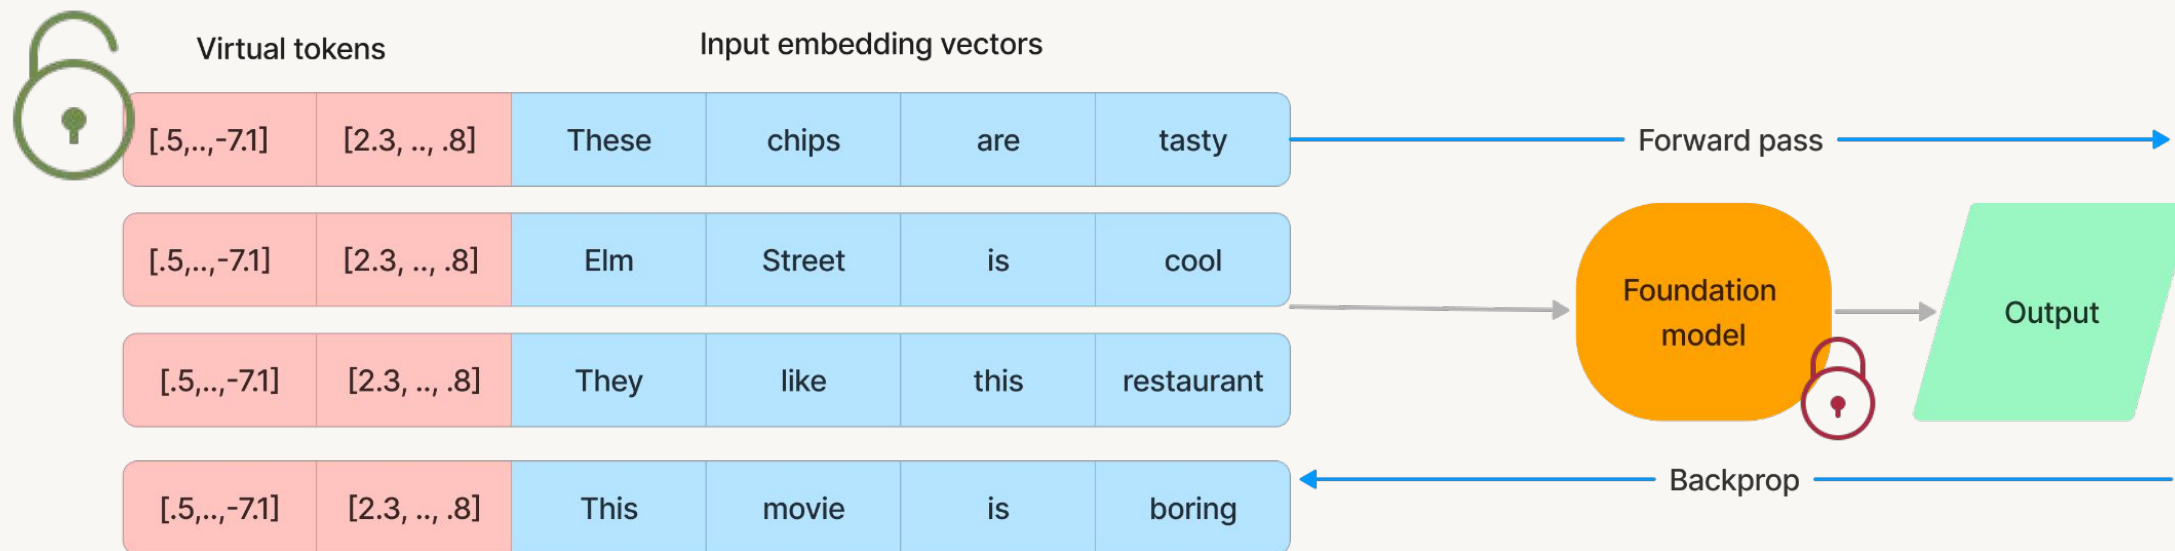
# Compare full fine-tuning vs prompt tuning

## Scenario: prompt tuning

Backprop: update **only prompt** weights based on loss

- The model learns the optimal representation of the prompt automatically

Task: Classify sentiment; task batch = 4

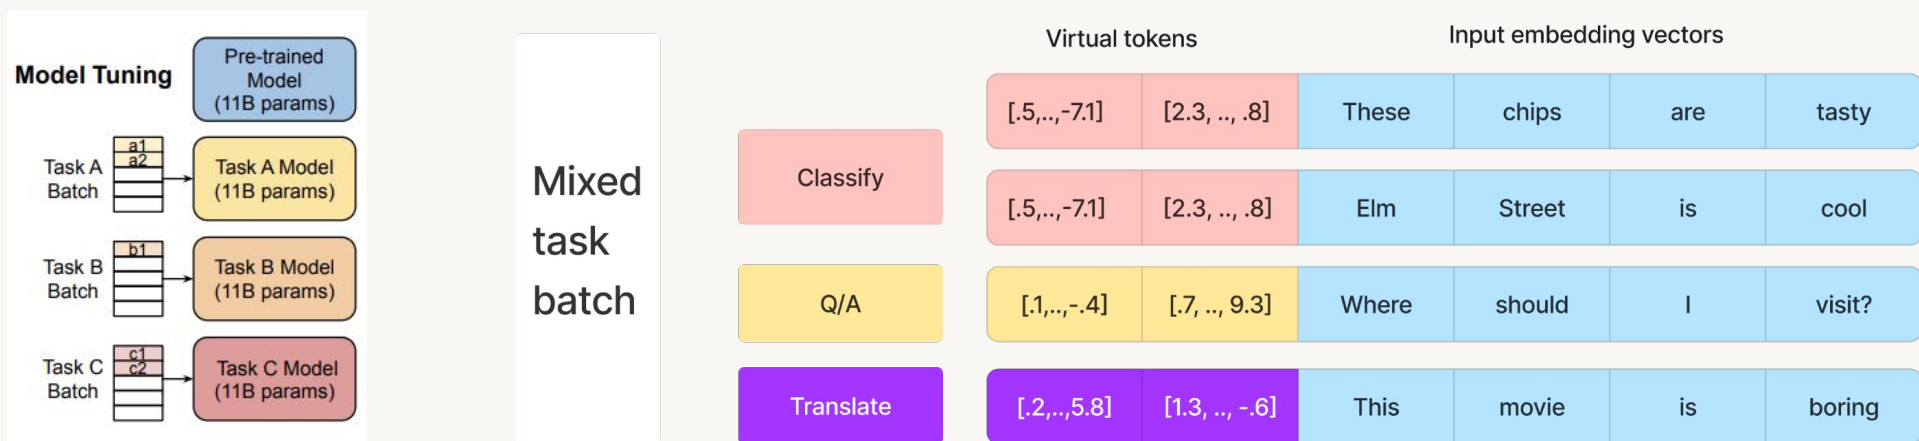




# Allows swapping of task prompts

Efficient for multi-task serving

- Each task is a prompt, not a model
  - Only need to serve a single copy of the frozen model for multi-task serving
- Prompts for various tasks can be applied to different inputs
  - A serving request can be a single, larger mixed task batch




Source: [Lester et al 2021](#)



# Matches fine tuning performance for >11B model

- Comparable with full fine-tuning at the 10B model scale
- More applicable to larger models
- SuperGLUE (2019)
  - Styled after GLUE, but more difficult and diverse
  - Boolean questions, comprehension, etc.

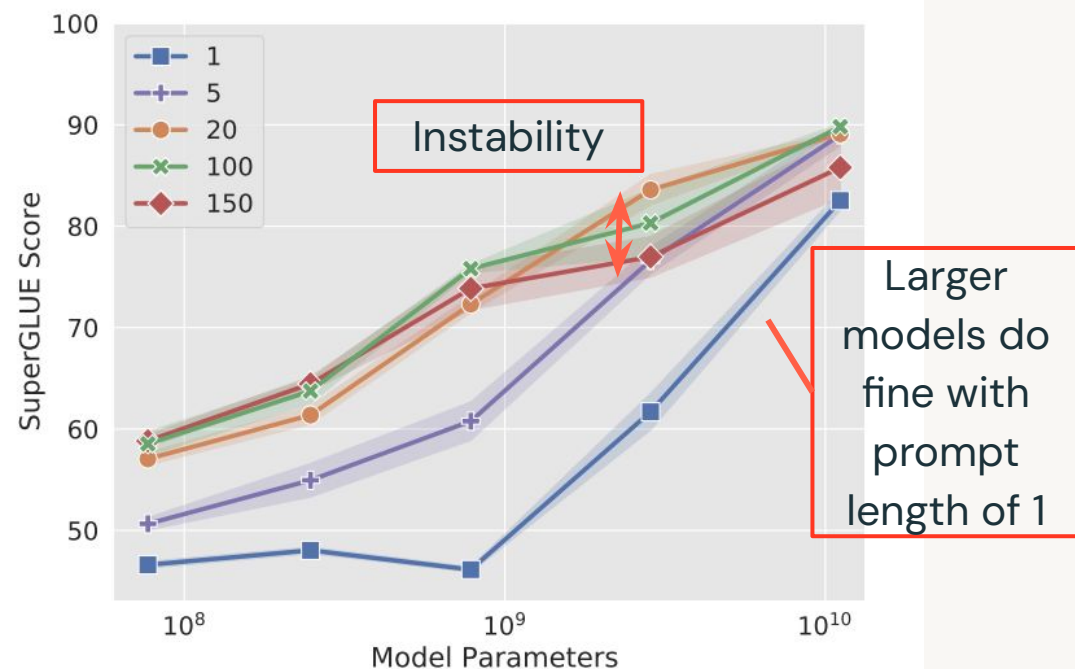


Source: [Lester et al 2021](#)



# Prompt length affects larger models less

Prompt length of 20–100 is typical



Source: [Lester et al 2021](#)

In this example:

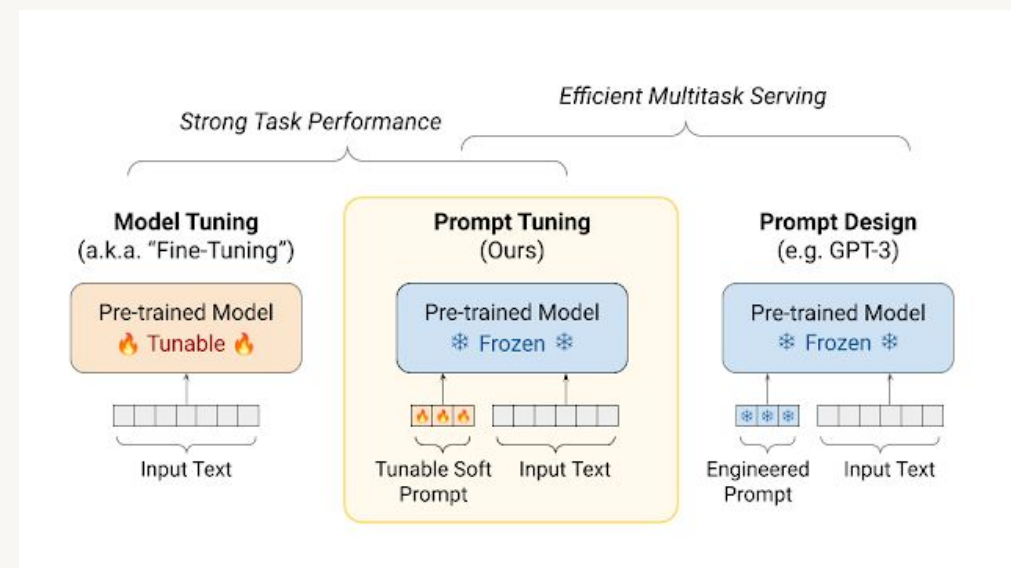
- (Virtual) prompt length = 2

| Virtual tokens |                | Input embedding vectors |        |     |        |
|----------------|----------------|-------------------------|--------|-----|--------|
| [.5,..,-7.1]   | [2.3, .., .8]  | These                   | chips  | are | tasty  |
| [.5,..,-7.1]   | [2.3, .., .8]  | Elm                     | Street | is  | cool   |
| [.1,..,-.4]    | [.7, .., 9.3]  | Where                   | should | I   | visit? |
| [.2,..,5.8]    | [1.3, .., -.6] | This                    | movie  | is  | boring |



# Advantages of prompt tuning

- Use whole training set
  - Not limited by # of examples that can fit in the context
- Automatically learn a new prompt for a new model
  - Backprop helps us find the best representation
- One foundation model copy only
- Resilient to domain shift



Source: [Google AI Blog](#)




# Disadvantages of prompt tuning

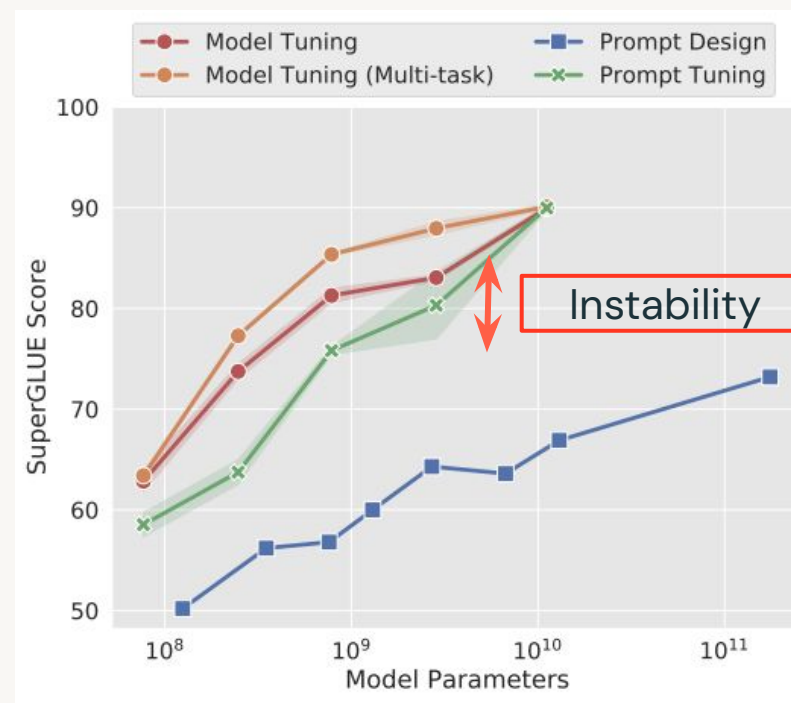
## Less interpretable

- Need to convert the embeddings back to tokens
- Use cosine distance to find the top-K nearest neighbors

Find which tokens are nearest to the virtual tokens



## Unstable performance

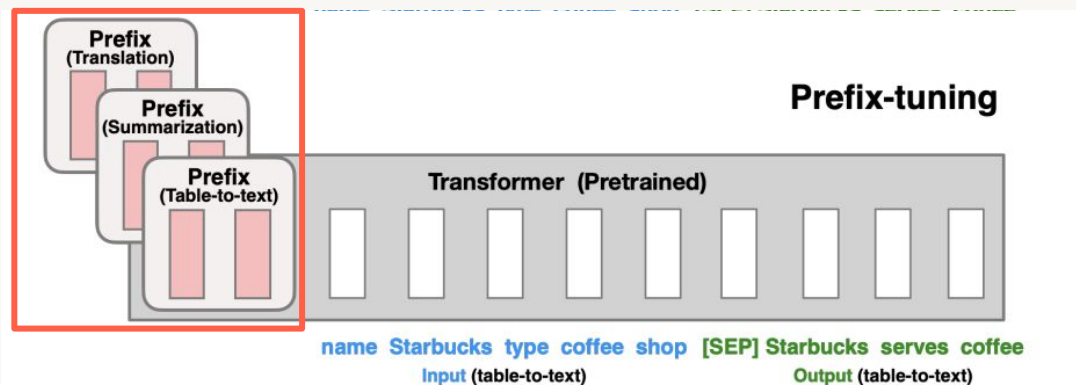


Source: [Lester et al 2021](#)

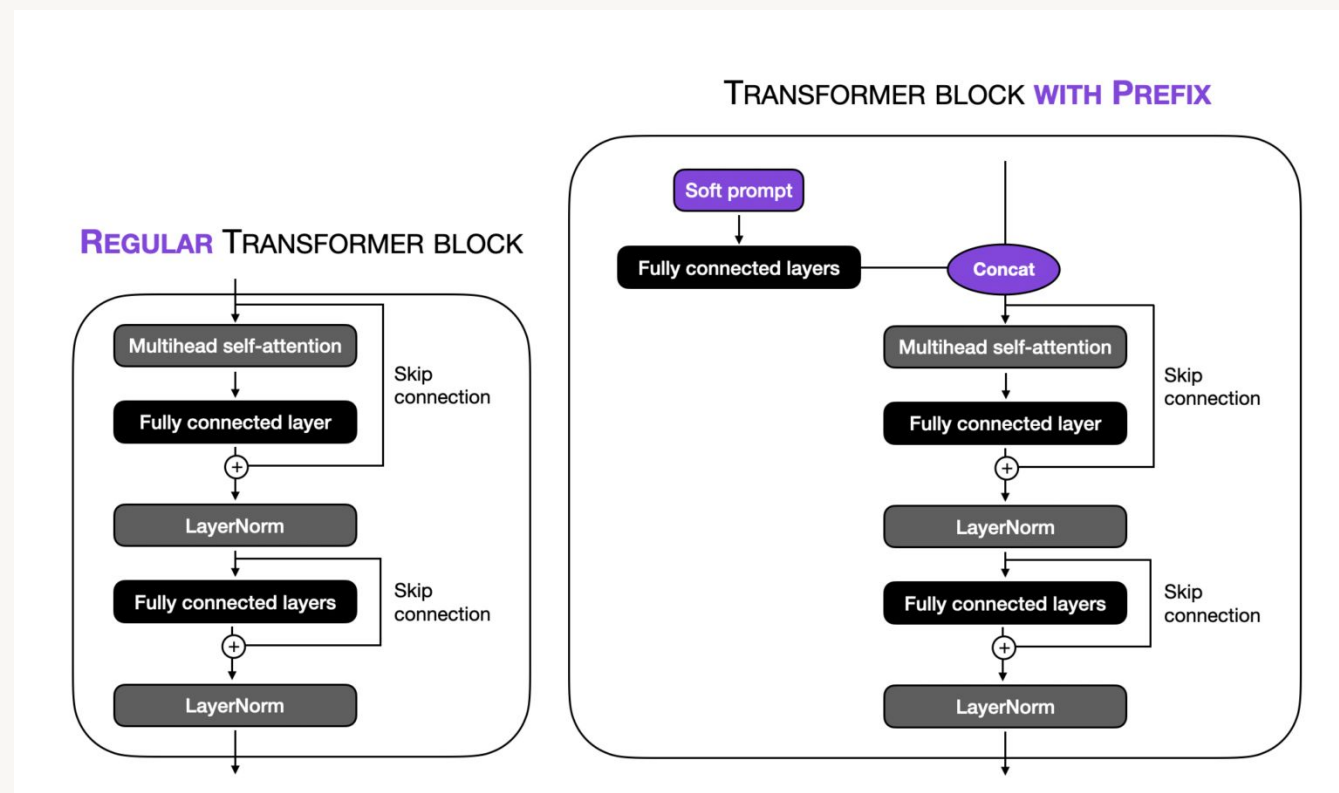


# Prefix tuning is very similar to prompt tuning

Adding tunable layer to each transformer block, rather than just the input layer



Source: [Li and Liang 2021](#)



Source: [Lightning AI](#)

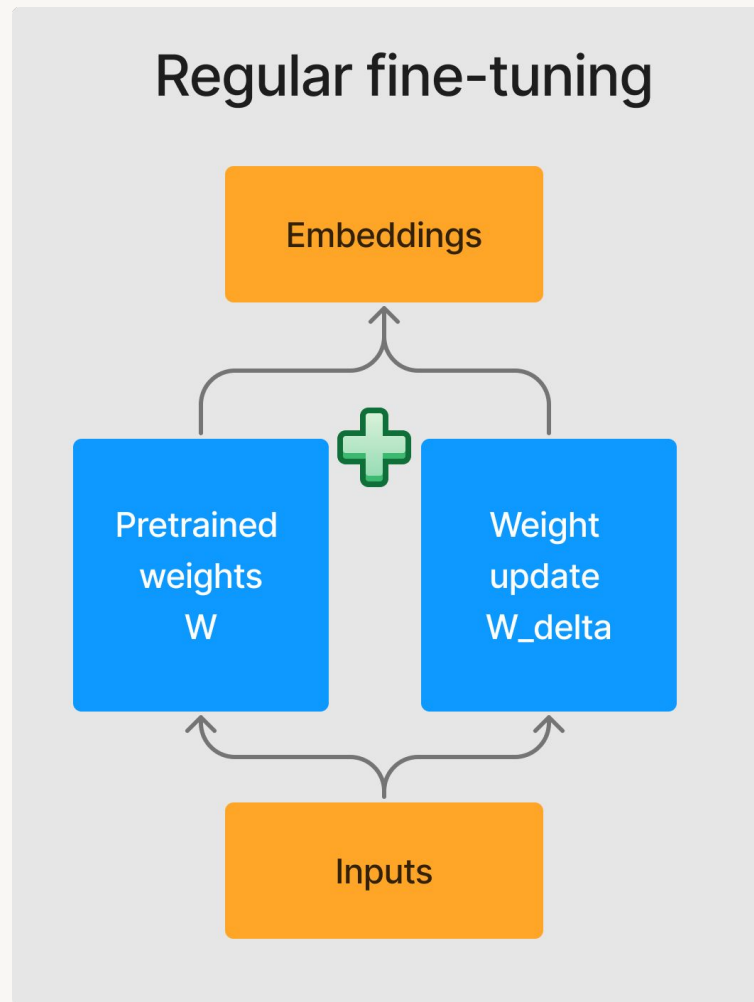


# Re-parameterization: LoRA



# Low-Rank Adaptation (LoRA)

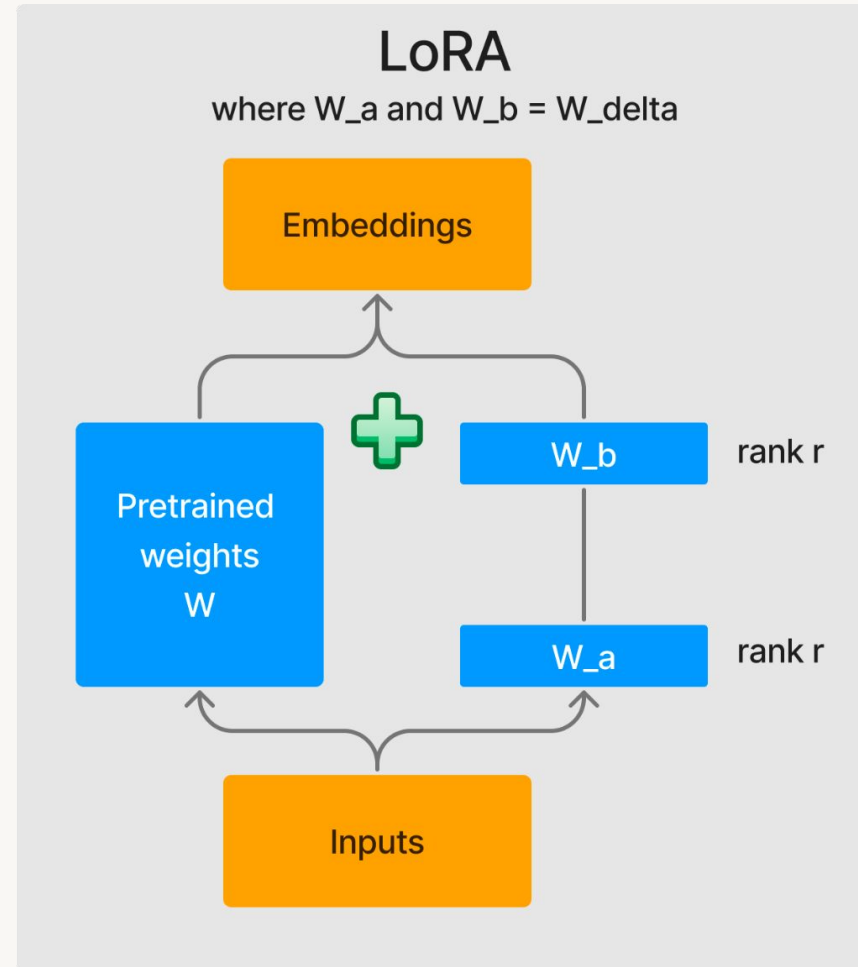
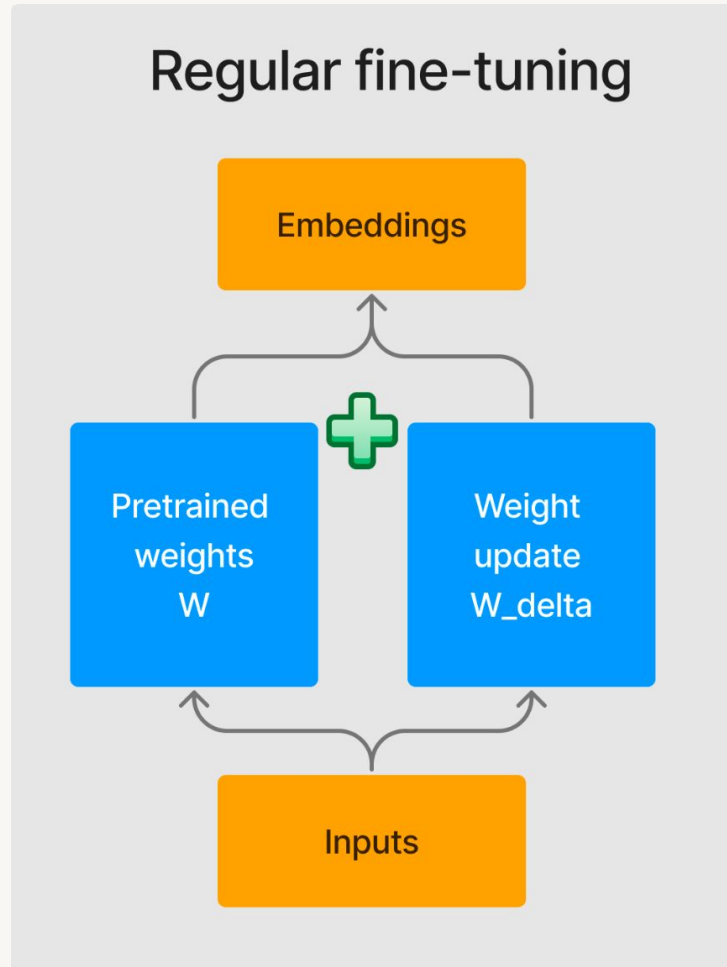
Decomposes the weight change matrix into lower-rank matrices





# Low-Rank Adaptation (LoRA)

Decomposes the weight change matrix into lower-rank matrices



# Rank? Brief visit to linear algebra

Maximum # of linearly independent columns or rows

- How many unique rows or columns?
- Full rank = no redundant row or column in the matrix
- Linear = can multiply by a constant
- Independence = no dependence on each other

$$\begin{bmatrix} 1 & 2 & 3 \\ 3 & 6 & 9 \end{bmatrix}$$

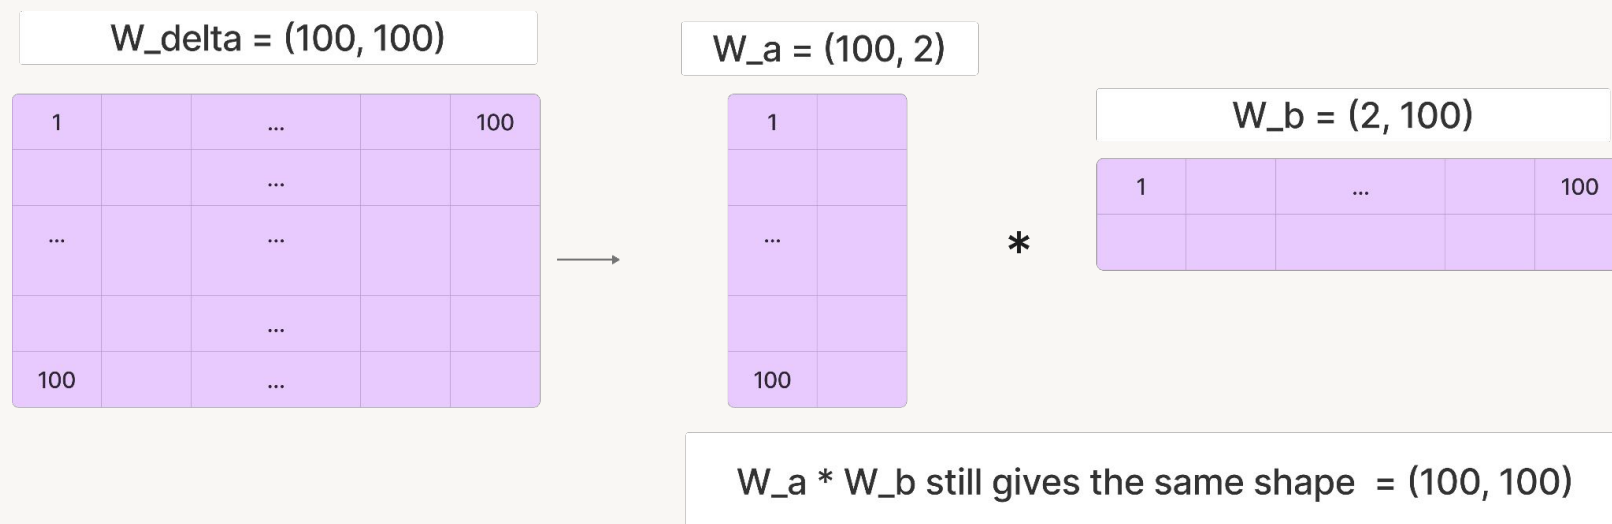
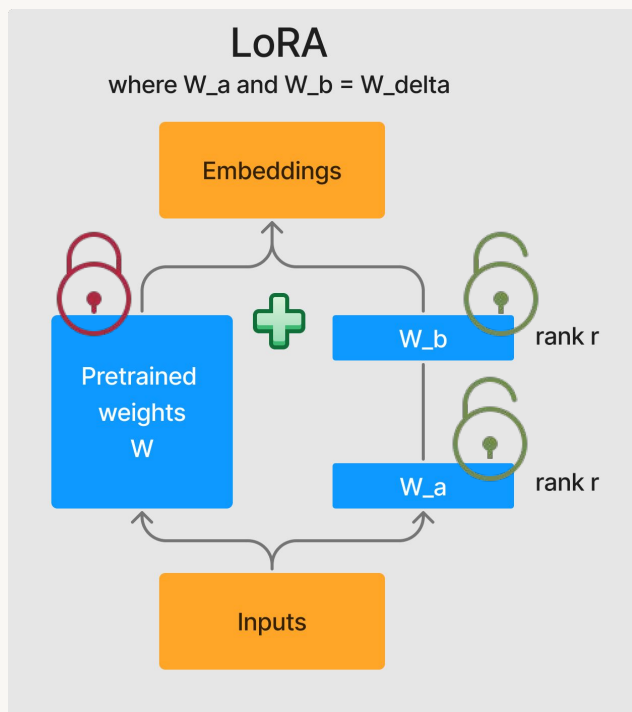
- Row rank: 1
  - 2nd row = 3x 1st row
- Column rank: 1
  - 2nd column = 2x 1st column
  - 3rd column = 2x 2nd column



# How does weight matrix decomposition work?

Observation: Actual rank of the attention weight matrices is low

$$W_{\text{delta}} = W_a * W_b$$



- Total parameters =  $(100 \times 2) + (2 \times 100) = 400$
- Original parameters =  $(100 \times 100) = 10,000$  parameters
- Reduction =  $10,000 - 400 = 96\%$ !



# LoRA matches/~outperforms full fine-tuning

- $37.7 / 175255.8$   
= 0.0002  
= 0.02% of parameters!

| Model&Method                  | # Trainable Parameters | WikiSQL     | MNLI-m      | SAMSum                | Rouge |
|-------------------------------|------------------------|-------------|-------------|-----------------------|-------|
|                               |                        | Acc. (%)    | Acc. (%)    | R1/R2/RL              |       |
| GPT-3 (FT)                    | 175,255.8M             | <b>73.8</b> | 89.5        | 52.0/28.0/44.5        |       |
| GPT-3 (BitFit)                | 14.2M                  | 71.3        | 91.0        | 51.3/27.4/43.5        |       |
| GPT-3 (PreEmbed)              | 3.2M                   | 63.1        | 88.6        | 48.3/24.2/40.5        |       |
| GPT-3 (PreLayer)              | 20.2M                  | 70.1        | 89.5        | 50.8/27.3/43.5        |       |
| GPT-3 (Adapter <sup>H</sup> ) | 7.1M                   | 71.9        | 89.8        | 53.0/28.9/44.8        |       |
| GPT-3 (Adapter <sup>H</sup> ) | 40.1M                  | 73.2        | <b>91.5</b> | 53.2/29.0/45.1        |       |
| GPT-3 (LoRA)                  | 4.7M                   | 73.4        | <b>91.7</b> | <b>53.8/29.8/45.9</b> |       |
| GPT-3 (LoRA)                  | 37.7M                  | <b>74.0</b> | <b>91.6</b> | 53.4/29.2/45.1        |       |

Source: [Hu et al 2021](#)



# LoRA performs well with very small ranks

GPT-3's validation accuracies are similar across rank sizes

$W_q$  = query

$W_k$  = key

$W_v$  = value

$W_o$  = output

|                          | Weight Type          | $r = 1$ | $r = 2$ | $r = 4$ | $r = 8$ | $r = 64$ |
|--------------------------|----------------------|---------|---------|---------|---------|----------|
| WikiSQL( $\pm 0.5\%$ )   | $W_q$                | 68.8    | 69.6    | 70.5    | 70.4    | 70.0     |
|                          | $W_q, W_v$           | 73.4    | 73.3    | 73.7    | 73.8    | 73.5     |
|                          | $W_q, W_k, W_v, W_o$ | 74.1    | 73.7    | 74.0    | 74.0    | 73.9     |
| MultiNLI ( $\pm 0.1\%$ ) | $W_q$                | 90.7    | 90.9    | 91.1    | 90.7    | 90.7     |
|                          | $W_q, W_v$           | 91.3    | 91.4    | 91.3    | 91.6    | 91.4     |
|                          | $W_q, W_k, W_v, W_o$ | 91.2    | 91.7    | 91.7    | 91.5    | 91.4     |

Source: [Hu et al 2021](#)

But, small  $r$  likely won't work for all tasks/datasets.

- E.g. downstream task is in a different language

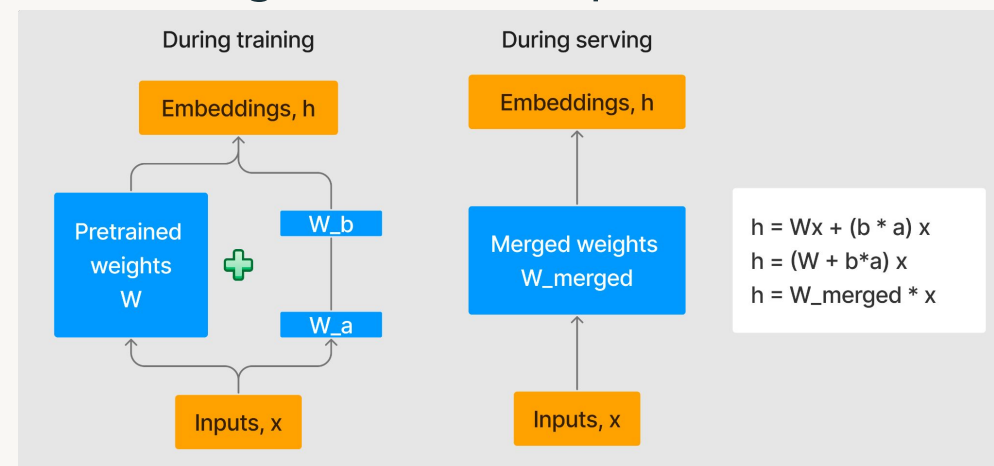


# Advantages of LoRA

Similar to prompt tuning, majority of the model weights are frozen

- Able to share and re-use the foundation model
  - Swap different LoRA weights for serving different tasks
- Improves training efficiency
  - Lower hardware barrier (no need to calculate most gradients or optimizer states)

- Adds no additional serving latency
  - $W_a * W_b$  can be merged



- Can be combined with other PEFT methods



# Limitations of LoRA

- Not straightforward to do multi-task serving
  - How to swap different combos of A and B in a **single** forward pass?
  - If dynamically choose A and B based on tasks, there is additional serving latency
- Future research
  - From LoRA authors: If  $W_{\text{delta}}$  is rank-deficient, is  $W$  too?
  - Newer PEFT technique: [IA3 \(2022\)](#)
    - Reduces even more trainable parameters than LoRA!



# PEFT Limitations





# Model performance limitations

- Difficult to match the performance of full fine-tuning
  - Sensitive to hyperparameters
  - Unstable performance
- Current research area: where is best to apply PEFT?
  - E.g. why apply PEFT to only attention weight matrices? Soft prompts?
  - [Vu et al 2022](#): Soft prompt transfer
- We may still need full-parameter fine-tuning
  - [Lv et al 2023](#) (released in June): use new optimizer, LOMO, to reduce memory usage to ~11%



# Compute limitations



Doesn't always make  
**inference** more efficient



Doesn't reduce the cost of  
**storing** massive foundation  
models



Doesn't reduce time  
complexity of **training**  
Requires full forward and  
backward passes



# Data Preparation Best Practices



# Better models from better training data

Many newer good models use C4 (e.g. MPT-7B)

## Llama

- Trained on 20 most-spoken languages, focusing on those with Latin and Cyrillic alphabets


Colossal  
Cleaned  
Crawled  
Corpus

| Dataset       | Sampling prop. | Epochs | Disk size |
|---------------|----------------|--------|-----------|
| CommonCrawl   | 67.0%          | 1.10   | 3.3 TB    |
| C4            | 15.0%          | 1.06   | 783 GB    |
| Github        | 4.5%           | 0.64   | 328 GB    |
| Wikipedia     | 4.5%           | 2.45   | 83 GB     |
| Books         | 4.5%           | 2.23   | 85 GB     |
| ArXiv         | 2.5%           | 1.06   | 92 GB     |
| StackExchange | 2.0%           | 1.03   | 78 GB     |

Source: [Touvron et al 2023](#)

## GPT-Neo and GPT-J

- Trained on the Pile: 22 diverse datasets
- Outperformed GPT-3 in some instances ([Read more here](#))



Source: [Gao et al 2020](#)



# Training data makes the biggest difference

Not necessarily the model architecture

- Bloomberg created 363B-token dataset of English financial documents spanning 40 years
  - Augmented with 345B-token public dataset
- Outperforms existing open models on financial tasks

| <i>Finance-Specific</i>                         | <b>BloombergGPT</b> | <b>GPT-NeoX</b> | <b>OPT-66B</b> | <b>BLOOM-176B</b> |
|---|---------------------|-----------------|----------------|-------------------|
| <b>Financial Tasks</b>                          | <b>62.51</b>        | 51.90           | 53.01          | 54.35             |
| <b>Bloomberg Tasks<br/>(Sentiment Analysis)</b> | <b>62.47</b>        | 29.23           | 35.76          | 33.39             |

| <i>General-Purpose</i>           | <b>BloombergGPT</b> | <b>GPT-NeoX</b> | <b>OPT-66B</b> | <b>BLOOM-176B</b> | <b>GPT-3</b> |
|----------------------------------|---------------------|-----------------|----------------|-------------------|--------------|
| <b>MMLU</b>                      | <b>39.18</b>        | 35.95           | 35.99          | 39.13             | <b>43.9</b>  |
| <b>Reading<br/>Comprehension</b> | <b>61.22</b>        | 42.81           | 50.21          | 49.37             | <b>67.0</b>  |
| <b>Linguistic Scenarios</b>      | <b>60.63</b>        | 57.18           | 58.59          | 58.26             | <b>63.4</b>  |

Source: [Wu et al 2023](#)



# How much fine-tuning data do I need?

- [Zhou et al 2023](#) (May): fine-tune 1,000 high-**quality** labeled examples from LLaMa 65B
  - When scaling up data **quantity**, need to scale up prompt **diversity**
- [OpenAI](#): At least a couple hundred
  - Doubling dataset size leads to linear increase in model performance
- How to get more data? Synthetic data
  - Synonym replacement / rewrite
  - Word deletion: “brilliantly expressed” => “expressed”
  - Word position swapping: “It is lovely” -> “Lovely, it is”
  - Noise injection: introduce typos



# Data preparation best practices

## Quantity, diversity, and quality

- Don't provide detailed instructions. Only prompt and completion.
  - Fixed separator `\n\n###\n\n` to inform when the prompt ends and completion begins
  - The separator shouldn't appear anywhere else

```
1 {"prompt": "<prompt text>", "completion": "<ideal generated text>"}  
2 {"prompt": "<prompt text>", "completion": "<ideal generated text>"}  
3 {"prompt": "<prompt text>", "completion": "<ideal generated text>"}  
4 ...
```

Source: [OpenAI](#)

### [Manually Authored – Writing]

**I am planning to start a book club with some friends. Can you write an email invitation for the same?**

Here is an example email invitation for a book club:

Subject: Join me for a book club!

Hi [Friend's Name], how are you?

I'm starting a book club with some friends, and I'd love to have you join us!

The book club is a casual and fun way to read and discuss books that we love or want to learn more about, and also make new friends. We plan to meet more-or-less once a month.

We're thinking of having our first meeting on [Date & Time]. Could that work for you?

Would love to see you there!

Thanks,  
[Your Name]

Source: [Zhou et al 2023](#)



# Data preparation best practices

- Remove undesired data
  - Offensive, toxic content
  - Private or confidential information
- Using LLM output as data is not always the answer
  - Imitation models learn style, rather than content ([Gudibande et al 2023](#))
  - Consistent with [Zhou et al 2023](#): knowledge is largely learned during pre-training
- Manually verify data quality





# Module Summary

## Efficient Fine-Tuning – What have we learned?

- Fine-tuning gives the best results, but can be computationally expensive
- Parameter-efficient fine-tuning reduces # of trainable parameters
- Prompt tuning allows virtual prompts to be learned automatically
- LoRA decomposes the weight change matrix into lower-rank matrices
- Fine-tuning data quality and diversity matters a lot



# Time for some code!



# Course Outline

Course Introduction

Module 1 – Transformers: Attention and the Transformer Architecture

Module 2 – Parameter Efficient Fine-Tuning: Doing more with less

Module 3 – Deployment Optimizations: Improving model size and speed

Module 4 – Multi-modal LLMs: Beyond text-based transformers

